

Session 23: More Kafka Assignment 1

TASK-1:

Create a java program MyKafkaProducer.java that takes a file name and delimiter as input arguments. It should read the content of file line by line.

Fields in the file are in following order

1. Kafka Topic Name
2. Key
3. value

For every line, insert the key and value to the repsective Kafka broker in a fire and forget mode.

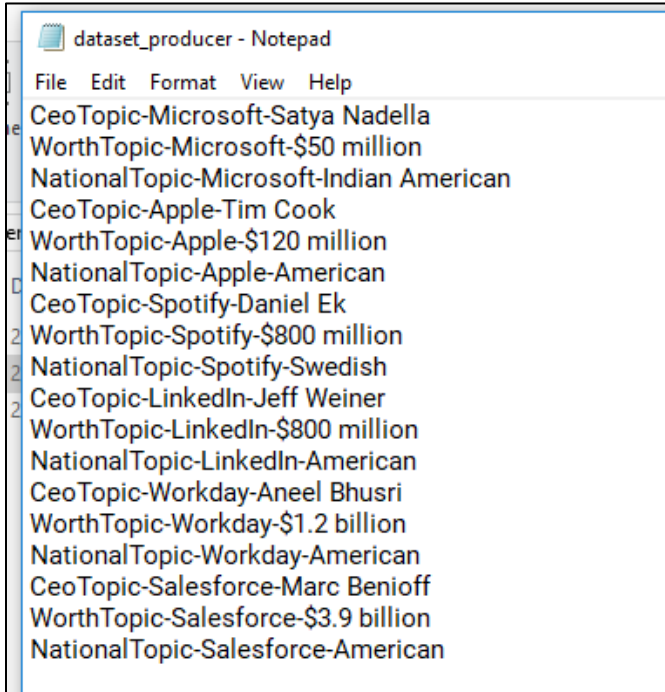
After record is sent, it should print appropriate message on screen.

Pass **dataset_producer.txt** as the input file and -as delimiter.

LINK: https://drive.google.com/file/d/0B_Qjau8wv1KoSnR5eHpKOF9rTFU/view?usp=sharing

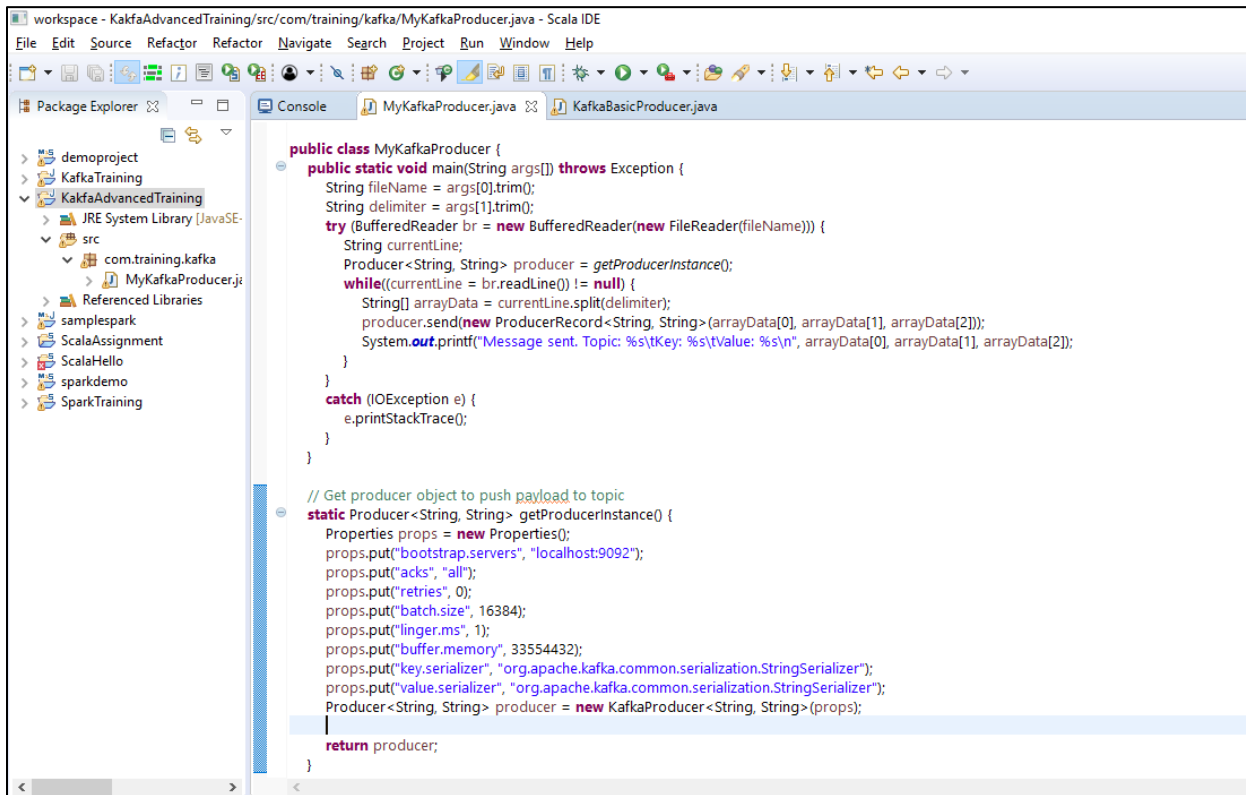
Note: The above input data file path (google drive) given in assignment PDF is invalid. No file exist in the path. So, I have created a sample file for my own.

Input: dataset_producer.txt



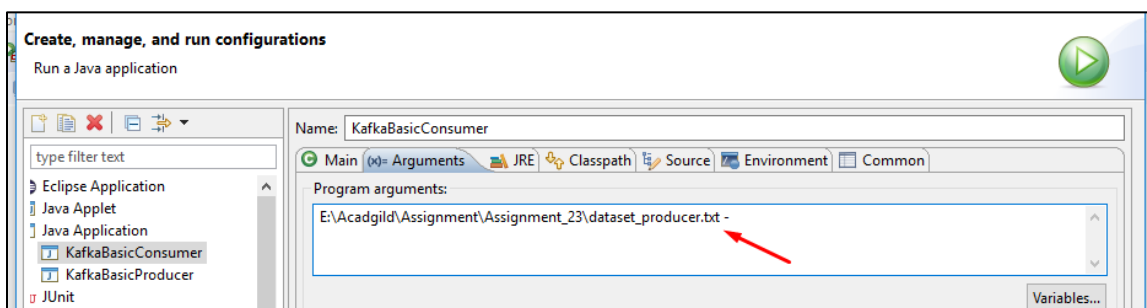
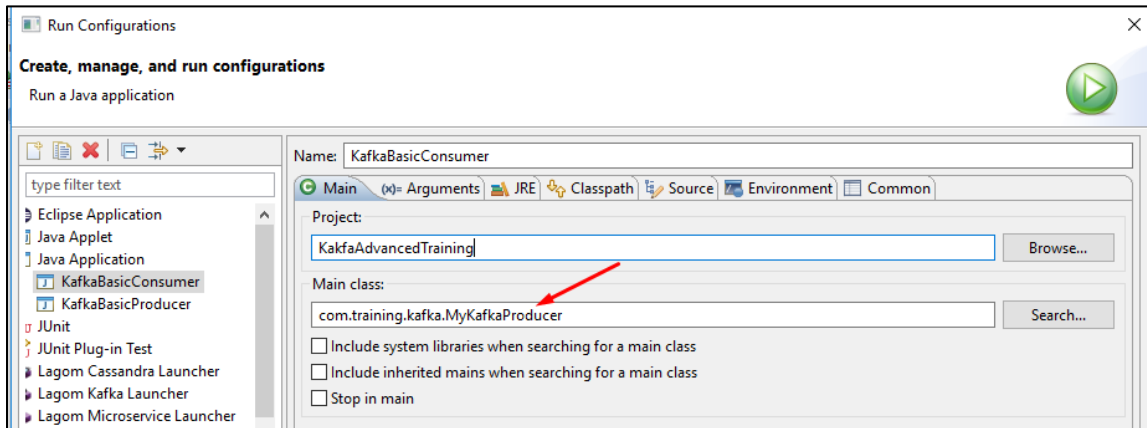
```
dataset_producer - Notepad
File Edit Format View Help
CeoTopic-Microsoft-Satya Nadella
WorthTopic-Microsoft-$50 million
NationalTopic-Microsoft-Indian American
CeoTopic-Apple-Tim Cook
WorthTopic-Apple-$120 million
NationalTopic-Apple-American
CeoTopic-Spotify-Daniel Ek
WorthTopic-Spotify-$800 million
NationalTopic-Spotify-Swedish
CeoTopic-LinkedIn-Jeff Weiner
WorthTopic-LinkedIn-$800 million
NationalTopic-LinkedIn-American
CeoTopic-Workday-Aneel Bhusri
WorthTopic-Workday-$1.2 billion
NationalTopic-Workday-American
CeoTopic-Salesforce-Marc Benioff
WorthTopic-Salesforce-$3.9 billion
NationalTopic-Salesforce-American
```

Program:



```
public class MyKafkaProducer {  
    public static void main(String args[]) throws Exception {  
        String fileName = args[0].trim();  
        String delimiter = args[1].trim();  
        try (BufferedReader br = new BufferedReader(new FileReader(fileName))) {  
            String currentLine;  
            Producer<String, String> producer = getProducerInstance();  
            while((currentLine = br.readLine()) != null) {  
                String[] arrayData = currentLine.split(delimiter);  
                producer.send(new ProducerRecord<String, String>(arrayData[0], arrayData[1], arrayData[2]));  
                System.out.printf("Message sent. Topic: %s\tKey: %s\tValue: %s\n", arrayData[0], arrayData[1], arrayData[2]);  
            }  
        } catch (IOException e) {  
            e.printStackTrace();  
        }  
    }  
  
    // Get producer object to push payload to topic  
    static Producer<String, String> getProducerInstance() {  
        Properties props = new Properties();  
        props.put("bootstrap.servers", "localhost:9092");  
        props.put("acks", "all");  
        props.put("retries", 0);  
        props.put("batch.size", 16384);  
        props.put("linger.ms", 1);  
        props.put("buffer.memory", 33554432);  
        props.put("key.serializer", "org.apache.kafka.common.serialization.StringSerializer");  
        props.put("value.serializer", "org.apache.kafka.common.serialization.StringSerializer");  
        Producer<String, String> producer = new KafkaProducer<String, String>(props);  
        return producer;  
    }  
}
```

Execution:



Execution Output:

A screenshot of a Java IDE's console window. The title bar shows 'Console', 'MyKafkaProducer.java', and 'KafkaBasicProducer.java'. The console output shows a series of messages sent to three topics: CeoTopic, WorthTopic, and NationalTopic. Each message includes a key and a value. The messages are grouped by topic, with CeoTopic having 4 messages, WorthTopic having 4 messages, and NationalTopic having 4 messages. The messages are: 1. CeoTopic (Key: Microsoft, Value: Satya Nadella), 2. WorthTopic (Key: Microsoft, Value: \$50 million), 3. NationalTopic (Key: Microsoft, Value: Indian American), 4. CeoTopic (Key: Apple, Value: Tim Cook), 5. WorthTopic (Key: Apple, Value: \$120 million), 6. NationalTopic (Key: Apple, Value: American), 7. CeoTopic (Key: Spotify, Value: Daniel Ek), 8. WorthTopic (Key: Spotify, Value: \$800 million), 9. NationalTopic (Key: Spotify, Value: Swedish), 10. CeoTopic (Key: LinkedIn, Value: Jeff Weiner), 11. WorthTopic (Key: LinkedIn, Value: \$800 million), 12. NationalTopic (Key: LinkedIn, Value: American), 13. CeoTopic (Key: Workday, Value: Aneel Bhusri), 14. WorthTopic (Key: Workday, Value: \$1.2 billion), 15. NationalTopic (Key: Workday, Value: American), 16. CeoTopic (Key: Salesforce, Value: Marc Benioff), 17. WorthTopic (Key: Salesforce, Value: \$3.9 billion), 18. NationalTopic (Key: Salesforce, Value: American).

```
<terminated> KafkaBasicConsumer [Java Application] C:\Program Files\Java\jre1.8.0_191\bin\javaw.exe (Feb 13, 2019, 5:23:38 PM)
Message sent. Topic: CeoTopic Key: Microsoft Value: Satya Nadella
Message sent. Topic: WorthTopic Key: Microsoft Value: $50 million
Message sent. Topic: NationalTopic Key: Microsoft Value: Indian American
Message sent. Topic: CeoTopic Key: Apple Value: Tim Cook
Message sent. Topic: WorthTopic Key: Apple Value: $120 million
Message sent. Topic: NationalTopic Key: Apple Value: American
Message sent. Topic: CeoTopic Key: Spotify Value: Daniel Ek
Message sent. Topic: WorthTopic Key: Spotify Value: $800 million
Message sent. Topic: NationalTopic Key: Spotify Value: Swedish
Message sent. Topic: CeoTopic Key: LinkedIn Value: Jeff Weiner
Message sent. Topic: WorthTopic Key: LinkedIn Value: $800 million
Message sent. Topic: NationalTopic Key: LinkedIn Value: American
Message sent. Topic: CeoTopic Key: Workday Value: Aneel Bhusri
Message sent. Topic: WorthTopic Key: Workday Value: $1.2 billion
Message sent. Topic: NationalTopic Key: Workday Value: American
Message sent. Topic: CeoTopic Key: Salesforce Value: Marc Benioff
Message sent. Topic: WorthTopic Key: Salesforce Value: $3.9 billion
Message sent. Topic: NationalTopic Key: Salesforce Value: American
```

As per the input file, there are 3 topics involved in publishing messages,

- CeoTopic
- WorthTopic
- NationalTopic

All 3 topics were subscribed and queried to see the messages. All the topics were displaying messages accordingly,

CeoTopic:

```
.\.bin\windows\kafka-console-consumer.bat --bootstrap-server localhost:9092 --topic CeoTopic --from-beginning
```

WorthTopic:

```
.\.bin\windows\kafka-console-consumer.bat --bootstrap-server localhost:9092 --topic WorthTopic --from-beginning
```

NationalTopic:

```
.\.bin\windows\kafka-console-consumer.bat --bootstrap-server localhost:9092 --topic NationalTopic --from-beginning
```

```
C:\Windows\system32\cmd.exe

E:\Hadoop\kafka_2.12-2.1.0>.\bin\windows\kafka-console-consumer.bat --bootstrap-server localhost:9092 --topic CeoTopic --from-beginning
Satya Nadella
Tim Cook
Daniel Ek
Jeff Weiner
Aneel Bhusri
Marc Benioff
Processed a total of 6 messages
Terminate batch job (Y/N)? y

E:\Hadoop\kafka_2.12-2.1.0>.\bin\windows\kafka-console-consumer.bat --bootstrap-server localhost:9092 --topic WorthTopic --from-beginning
$50 million
$120 million
$800 million
$800 million
$1.2 billion
$3.9 billion
Processed a total of 6 messages
Terminate batch job (Y/N)? y

E:\Hadoop\kafka_2.12-2.1.0>.\bin\windows\kafka-console-consumer.bat --bootstrap-server localhost:9092 --topic NationalTopic --from-beginning
Indian American
American
Swedish
American
American
Processed a total of 6 messages
Terminate batch job (Y/N)? y

E:\Hadoop\kafka_2.12-2.1.0>_
```

TASK-2:

Modify the previous program MyKafkaProducer.java and create a new Java program

KafkaProducerWithAck.java

This should perform the same task as of KafkaProducer.java with some modification.

When passing any data to a topic, it should wait for acknowledgement.

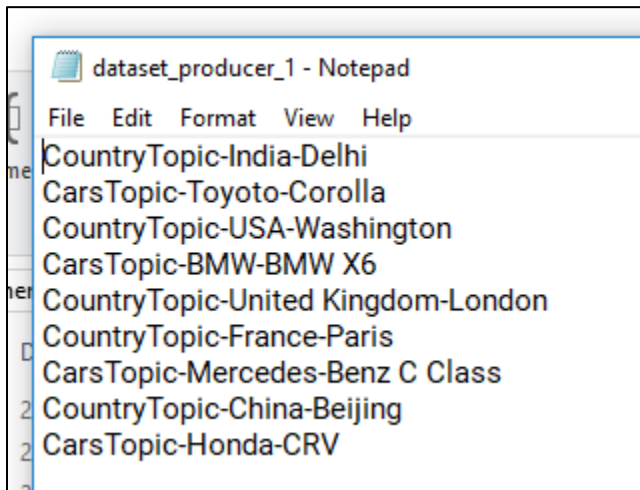
After acknowledgement is received from the broker, it should print the key and value which has been written to a specified topic.

The application should attempt for 3 retries before giving any exception.

Pass **dataset_producer.txt** as the input file and -as delimiter.

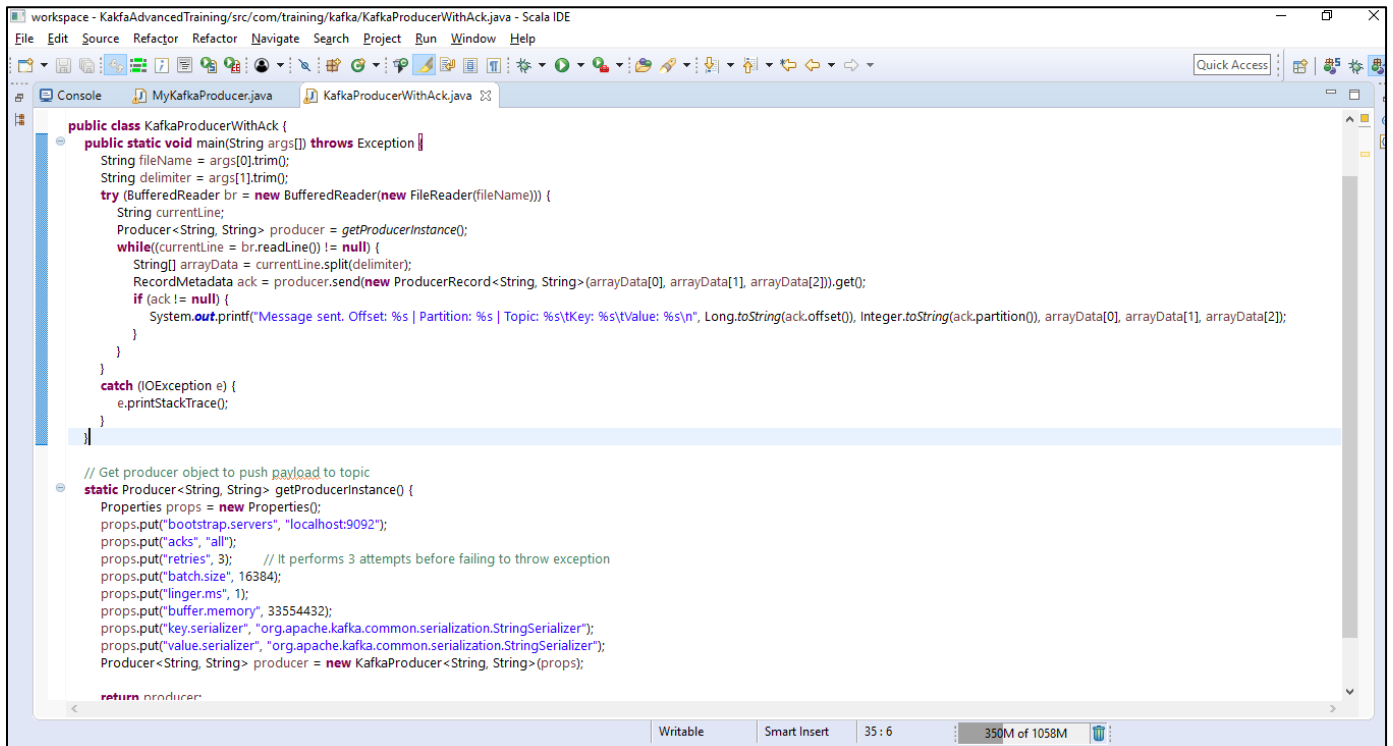
Note: The input data file path (google drive) given in assignment PDF is invalid. No file exist in the path. So, I have created a sample file for my own.

Input: dataset_producer_1.txt



```
dataset_producer_1 - Notepad
File Edit Format View Help
CountryTopic-India-Delhi
CarsTopic-Toyota-Corolla
CountryTopic-USA-Washington
CarsTopic-BMW-BMW X6
CountryTopic-United Kingdom-London
CountryTopic-France-Paris
CarsTopic-Mercedes-Benz C Class
CountryTopic-China-Beijing
CarsTopic-Honda-CRV
```

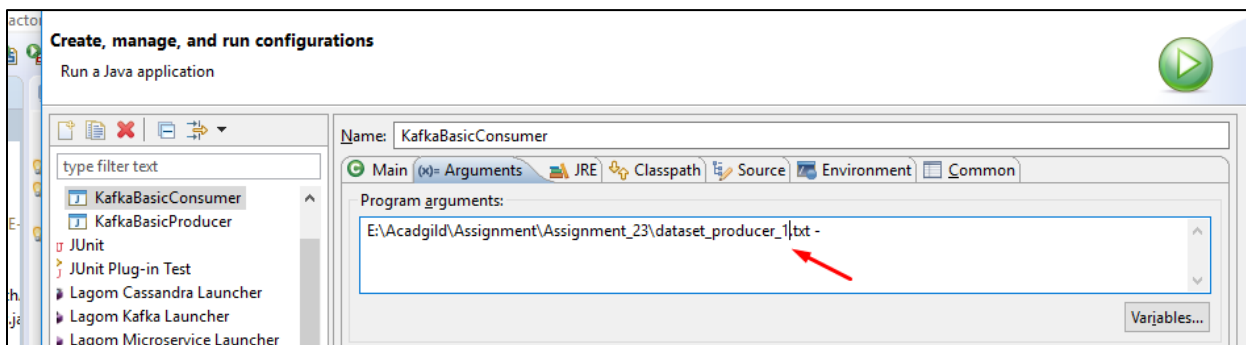
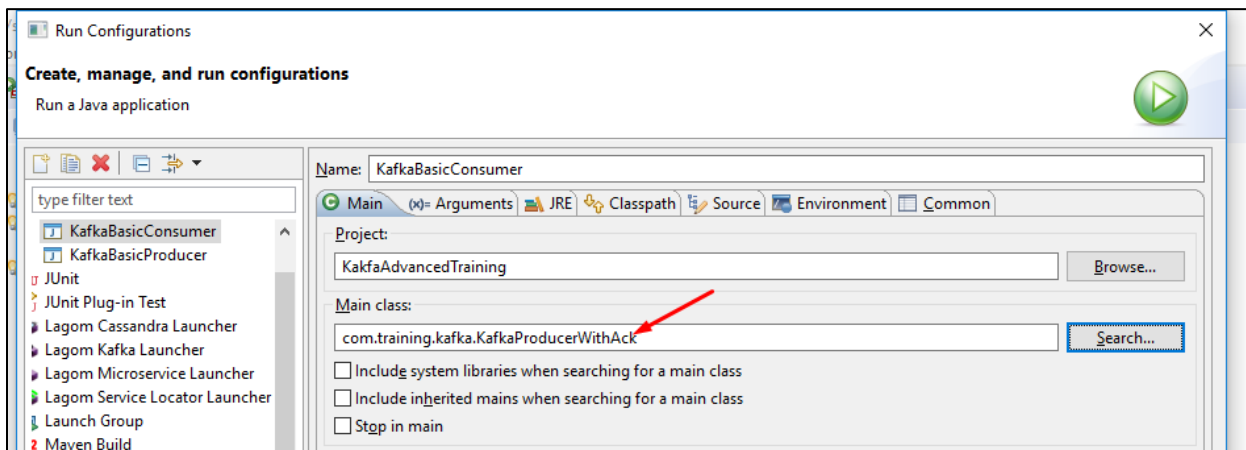
Program:



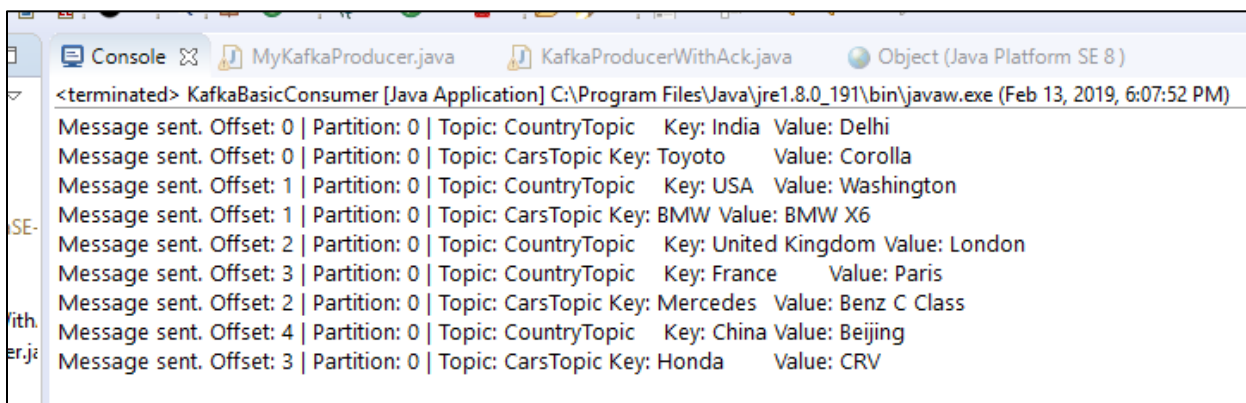
```
workspace - KafkaAdvancedTraining/src/com/training/kafka/KafkaProducerWithAck.java - Scala IDE
File Edit Source Refactor Refactor Navigate Search Project Run Window Help
Console MyKafkaProducer.java KafkaProducerWithAck.java
public class KafkaProducerWithAck {
    public static void main(String args[]) throws Exception {
        String fileName = args[0].trim();
        String delimiter = args[1].trim();
        try (BufferedReader br = new BufferedReader(new FileReader(fileName))) {
            String currentLine;
            Producer<String, String> producer = getProducerInstance();
            while((currentLine = br.readLine()) != null) {
                String[] arrayData = currentLine.split(delimiter);
                RecordMetadata ack = producer.send(new ProducerRecord<String, String>(arrayData[0], arrayData[1], arrayData[2])).get();
                if (ack != null) {
                    System.out.printf("Message sent. Offset: %s | Partition: %s | Topic: %s\tKey: %s\tValue: %s\n", Long.toString(ack.offset()), Integer.toString(ack.partition()), arrayData[0], arrayData[1], arrayData[2]);
                }
            }
        } catch (IOException e) {
            e.printStackTrace();
        }
    }

    // Get producer object to push payload to topic
    static Producer<String, String> getProducerInstance() {
        Properties props = new Properties();
        props.put("bootstrap.servers", "localhost:9092");
        props.put("acks", "all");
        props.put("retries", 3); // It performs 3 attempts before failing to throw exception
        props.put("batch.size", 16384);
        props.put("linger.ms", 1);
        props.put("buffer.memory", 33554432);
        props.put("key.serializer", "org.apache.kafka.common.serialization.StringSerializer");
        props.put("value.serializer", "org.apache.kafka.common.serialization.StringSerializer");
        Producer<String, String> producer = new KafkaProducer<String, String>(props);
        return producer;
    }
}
```

Console application argument settings:

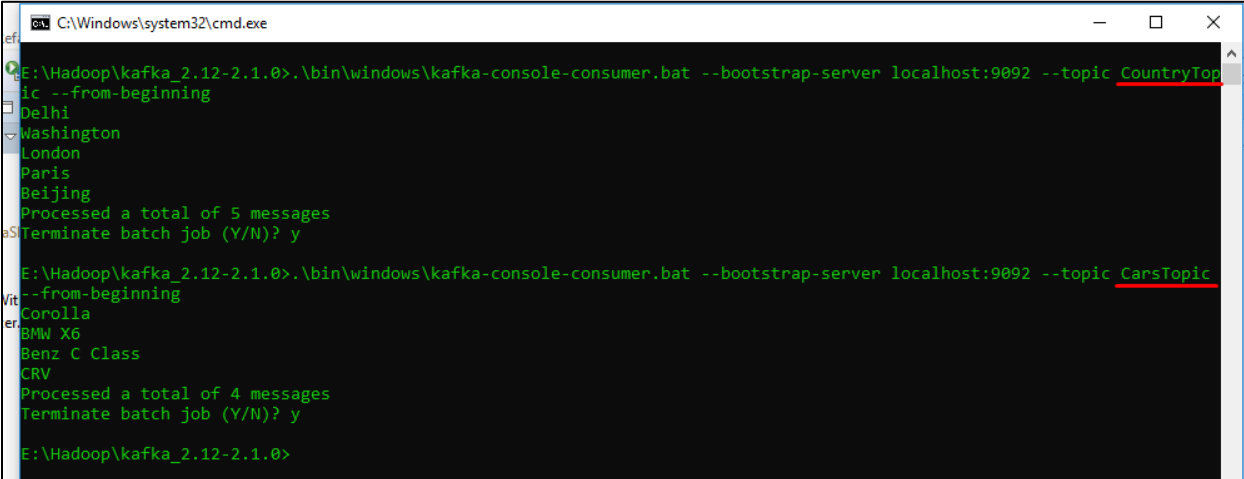


Execution Result:



Final Result:

Both topics (CountryTopic and CarsTopic) are subscribed and queried. The published messages are listed in it.



```
C:\Windows\system32\cmd.exe

E:\Hadoop\kafka_2.12-2.1.0>. \bin\windows\kafka-console-consumer.bat --bootstrap-server localhost:9092 --topic CountryTopic --from-beginning
Delhi
Washington
London
Paris
Beijing
Processed a total of 5 messages
Terminate batch job (Y/N)? y

E:\Hadoop\kafka_2.12-2.1.0>. \bin\windows\kafka-console-consumer.bat --bootstrap-server localhost:9092 --topic CarsTopic --from-beginning
Corolla
BMW X6
Benz C Class
CRV
Processed a total of 4 messages
Terminate batch job (Y/N)? y

E:\Hadoop\kafka_2.12-2.1.0>
```